F5. Per-question paired outcome, M5 SLM vs M4 SAGE-score greedy — TEST-1704 (n = 1,704), B = 4, maximum budget

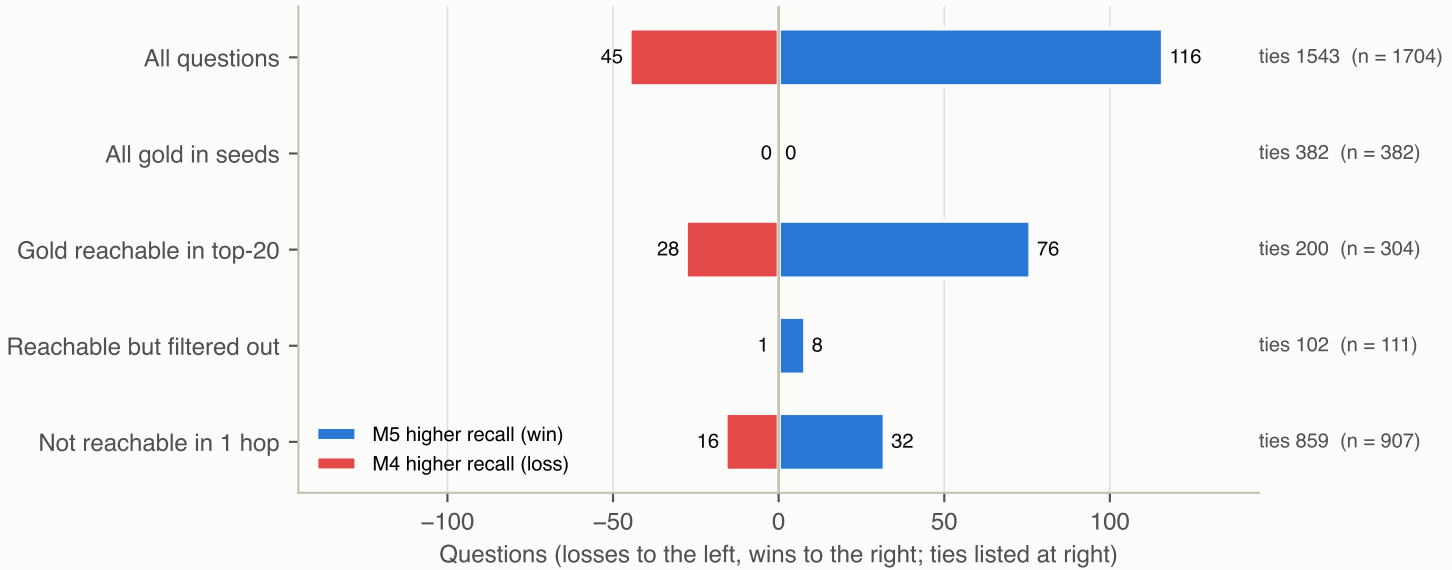

Win/tie/loss on per-question ID Recall (the T1 comparison: Δ = +0.0402, 95% CI [+0.0275, +0.0528], Holm p <0.0001). Categories as in T5c. Source: analysis/scores/per_question.jsonl, analysis/stats/comparisons.json.